ANONYMIZED PUBLIC SAMPLE

AI Engineer Candidate

United States · Contact details withheld

Candidate contact details and employer/institution names are masked. Role titles, dates, skills, experience descriptions, metrics, and project descriptions are retained from the supplied resume.

SKILLS

Programming & Data: Python, SQL, RESTful API Design, Microservices, Terraform
AI & MLOps: LLMs, BERT, Model Observability, Feature Engineering, Scikit-learn, PyTorch, Model Lifecycle
Cloud & Scale: AWS (Lambda, ECS, S3, API Gateway), Azure (AKS, Azure DevOps), Docker, Kubernetes (K8s)
Messaging & Streaming: RabbitMQ, Kafka, Spark Streaming, Redis, ActiveMQ
Process & Tools: CI/CD Pipelines, GitHub Actions, GitLab, Agile/Scrum, FHIR Interoperability

EXPERIENCE

Company A United States · Jan 2026 - Present

AI Engineer

- Architected FocusPal, an AI-driven productivity suite, utilizing AWS SageMaker Autopilot to deploy binary classification endpoints for real-time user distraction detection.
- Developed serverless data pipelines using AWS Lambda to ingest browser activity logs, achieving sub-second latency for inference results via high-performance REST endpoints.
- Engineered a robust activity logging system using OpenSearch (Elasticsearch service) with Basic Auth security, enabling real-time indexing and visualization of thousands of work sessions.
- Implemented automated data export workflows to Amazon S3 using Python and Boto3, creating labeled CSV datasets for continuous model retraining and drift analysis.

Company B India · Jul 2022 - Jul 2024

Software Engineer Associate

- Architected FHIR-compliant data exchange layers using Python and .NET Core, enabling seamless interoperability for 50+ enterprise healthcare providers while maintaining HIPAA compliance standards.
- Eliminated system bottlenecks by replacing legacy synchronous calls with a RabbitMQ-driven event-based architecture, resulting in a 12% increase in message throughput and 99.9% system availability.
- Engineered a production-ready ML pipeline using BERT-based Transformers to parse unstructured medical claims, reducing manual data-entry overhead by 14% and increasing extraction precision by 15%.
- Orchestrated a full migration of 15+ microservices to Azure AKS (Kubernetes), utilizing Docker for isolation and reducing deployment-related downtime by 30% across five enterprise projects.
- Developed automated CI/CD pipelines using Azure DevOps and GitLab, reducing software release cycles by 40% through integrated unit testing and automated security scans.
- Designed and optimized highly available RESTful APIs using ASP.NET Core and Spring Boot, serving millions of requests monthly for Medicaid-related financial platforms.
- Collaborated with cross-functional Agile teams to integrate third-party CRM platforms, ensuring real-time data synchronization and reducing document search latency by 40%.
- Authored Terraform scripts to manage Infrastructure as Code (IaC) on Azure, ensuring reproducible environments and reducing manual infrastructure setup time by 60%.

Research Institution A United States · Aug 2024 - Dec 2025

Graduate Research Assistant (AI/ML Engineering)

- Leading the technical design of MLOps workshops for 50+ graduate students, focusing on reproducible AI research, model drift monitoring, and automated hyperparameter tuning.
- Implemented automated testing suites for ML models (Deep Checks/Pytest), ensuring data integrity and model performance consistency across high-dimensional healthcare datasets.
- Developed scalable data ingestion pipelines using Pandas and Scikit-learn, implementing SMOTE techniques to handle class imbalance in credit risk and patient diagnostic models.
- Standardized model versioning and tracking using Git, enabling collaborative research on neural network architectures and NLP fine-tuning for 10+ student-led initiatives.

PROJECTS

- **Healthcare Sentiment Stream:** Built a real-time data ingestion engine using Kafka and Spark Streaming; implemented a custom NLP classifier to monitor Medicaid-related sentiment trends with 90% accuracy and sub-second latency.
- **BERT-Powered Spam Guard:** Developed a transformer-based classification engine for large-scale text streams; utilized advanced tokenization and LSTM hybrid architectures to achieve an 89% recall rate on adversarial samples.
- **Credit Risk Mitigation Engine:** Architected a binary classification model using Scikit-learn for credit risk assessment; optimized model hyperparameters to achieve a 12% improvement in F1-score over baseline models.

EDUCATION

- **University A** · United States

Master of Science in Data Science and AI; GPA: 3.6 · Jan 2024 - Dec 2025

- **University B** · India

Bachelor of Technology in Computer Engineering; GPA: 3.2 · Aug 2018 - Aug 2022